AI Job Search Assistant

Install & Run Guide - find jobs, keep your tracker, and tailor a resume for every role. Open source, runs on your own machine, with your own resume and your own AI key.

Code: https://github.com/ananyajoshi0203-cpu/ai-job-search-assistant

Live demo (sample data): https://muse.ai/s/ai-job-search-assistant-demo-gxz6wxtulxcgixs/-5S4SsvftIjOhdFyLUiDSA2IJZ6bmLaTIDBgBvq_Egs

1. Install

You need Python 3.10 or newer (check: python3 --version; if missing, install from python.org/downloads) and git.

```
git clone https://github.com/ananyajoshi0203-cpu/ai-job-search-assistant.git
cd ai-job-search-assistant
python3 -m venv .venv
source .venv/bin/activate        # Windows: .venv\Scripts\activate
python -m pip install -e .
```

Always use 'python -m pip', never bare 'pip' (on macOS, bare pip often does not exist). Every new terminal window needs the activate line again.

2. Run it

Easiest: double-click run.command (macOS), or run one command from the repo folder with the venv active:

```
jobhunt app
```

Your browser opens the app at http://localhost:8765. Everything is buttons:

- Run crawl - search the sources, dedupe against your tracker, verify each posting is still open
- Tailor Found roles - build the per-role resumes for roles already found
- All roles - your full tracker with status filters (Found / Ready to review / Applied / Skipped / Closed)
- Export .xlsx - download the tracker as a spreadsheet
- Each ready role is a card: click the resume name to read the tailored resume and approve it; Apply pre-fills the application form; No marks it Skipped forever

Command-line equivalents exist (jobhunt crawl, tailor, list, review, export). All commands share the same --config (default config.yaml) and --db (default data/tracker.db) - use the same pair everywhere.

3. Make it yours

```
cp config.example.yaml config.yaml
```

Then edit config.yaml:

- profile.name, profile.email_for_applications - used in your applications
- profile.roles, profile.locations, profile.keywords - what to search for
- profile.master_resume_path - point this at YOUR master resume (a Markdown file). The tool never edits this file; every tailored resume is a new file
- sources.greenhouse / lever / ashby - company board tokens to watch (the token is the name in the board URL, e.g. boards.greenhouse.io/gleanwork). LinkedIn, Wellfound and Jobright adapters are on by default - note they restrict automated access in their terms; the adapters use public pages only, rate-limited
- llm.model / llm.base_url - your AI provider (any OpenAI-compatible endpoint). Set your key in the

terminal, never in the file:

```
export OPENAI_API_KEY=your-key
```

Tip: try everything first with the built-in fictional sample profile (sample/sample_config.yaml) before using your real resume.

4. What it does - the heavy lifting

Finding jobs: it crawls your sources on every run, dedupes against your SQLite tracker (a role you mark Skipped is never shown again), checks each new posting is still open (closed ones are recorded as Closed and never tailored), and keeps every lead and status in one place - the tracker is the single source of truth, and all data persists on your disk across restarts.

Building resumes: for each role it writes a fresh, one-page version of your resume that rewords your own experience in the job description's language, reorders sections and bullets so the most relevant work comes first, and surfaces the required keywords - but only where your master resume already shows that skill. Shared keywords are listed on each review card so you can see exactly what matched.

Strictly no reinventing: a tailored resume may not contain any skill, tool, project, employer, or number that is not in your master resume. This is enforced by an automatic check - if a draft adds anything new, it is rejected. Draft application answers are written from the tailored resume only; fields your resume does not cover are left blank, never invented.

5. Troubleshooting quick hits

- 'pip: command not found' - use python -m pip (inside the venv) or python3 -m pip
- 'command not found: jobhunt' - activate the venv; the command is one word, jobhunt
- 'FileNotFoundError: config.yaml' - you skipped the cp step, or your --config/--db flags do not match the ones you crawled with
- 'invalid choice' for a command - your clone is behind: git pull
- Roles stuck at Found - they are found but not tailored yet: press 'Tailor Found roles' (needs your API key set)

6. What's still in progress

Last thing, honestly: you apply yourself. Full apply-automation is still in progress - the tool's Apply button pre-fills the form on supported sites and stops before Submit, but the final application is always yours to review and submit. What the tool does today is the heavy lifting around that step: maintaining your status and leads, and the resume tweaking for every role.